

AutoData

November 20, 2025

1 Simple Linear Regression on the Auto Dataset

Chapter 3 – Question 8 (Applied)

This analysis investigates the relationship between **horsepower** (predictor) and **mpg** (response) using a simple linear regression model. We will perform regression using `sm.OLS()` from `statsmodels`, evaluate the model, create plots, and interpret the results.

```
[7]: import pandas as pd
import statsmodels.api as sm
import matplotlib.pyplot as plt

auto = pd.read_csv("/home/mlahkim15/ve/Auto/Auto.csv")
auto.head()
```

```
[7]:    mpg  cylinders  displacement horsepower  weight  acceleration  year \
0   18.0          8         307.0       130     3504        12.0      70
1   15.0          8         350.0       165     3693        11.5      70
2   18.0          8         318.0       150     3436        11.0      70
3   16.0          8         304.0       150     3433        12.0      70
4   17.0          8         302.0       140     3449        10.5      70

      origin           name
0       1  chevrolet chevelle malibu
1       1           buick skylark 320
2       1  plymouth satellite
3       1           amc rebel sst
4       1           ford torino
```

1.1 Part (a) — Fit the Linear Regression Model

We model `mpg` as the response variable and `horsepower` as the predictor variable.

```
[8]: # Convert horsepower to numeric, coerce errors to NaN
auto['horsepower'] = pd.to_numeric(auto['horsepower'], errors='coerce')

# Drop rows with missing values in mpg or horsepower
auto = auto.dropna(subset=['horsepower', 'mpg'])
```

```
# Check types and first rows
print(auto.dtypes)
auto.head

X = auto["horsepower"]
y = auto["mpg"]

X = sm.add_constant(X)
model = sm.OLS(y, X).fit()
model.summary()
```

```
mpg           float64
cylinders      int64
displacement   float64
horsepower     float64
weight          int64
acceleration   float64
year            int64
origin          int64
name            object
dtype: object
```

[8]:

Dep. Variable:	mpg	R-squared:	0.606			
Model:	OLS	Adj. R-squared:	0.605			
Method:	Least Squares	F-statistic:	599.7			
Date:	Thu, 20 Nov 2025	Prob (F-statistic):	7.03e-81			
Time:	13:36:18	Log-Likelihood:	-1178.7			
No. Observations:	392	AIC:	2361.			
Df Residuals:	390	BIC:	2369.			
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	39.9359	0.717	55.660	0.000	38.525	41.347
horsepower	-0.1578	0.006	-24.489	0.000	-0.171	-0.145
Omnibus:	16.432	Durbin-Watson:	0.920			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	17.305			
Skew:	0.492	Prob(JB):	0.000175			
Kurtosis:	3.299	Cond. No.	322.			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

1.1.1 Interpretation of Regression Output

i. Is there a relationship between horsepower and mpg?

Yes — the p-value for horsepower is extremely small (much less than 0.05), which means the relationship is statistically significant.

ii. **How strong is the relationship?**

The R-squared value is around **0.60**, meaning about **60%** of the variation in mpg is explained by horsepower.

iii. **Is the relationship positive or negative?**

The coefficient for horsepower is **negative**, meaning as horsepower increases, mpg decreases.

iv. **Prediction for horsepower = 98**


We will calculate the predicted mpg and 95% confidence and prediction intervals below.

```
[9]: new_value = pd.DataFrame({"const": [1], "horsepower": [98]})  
model.get_prediction(new_value).summary_frame(alpha=0.05)
```

```
[9]:      mean   mean_se  mean_ci_lower  mean_ci_upper  obs_ci_lower  \  
0  24.467077  0.251262     23.973079     24.961075    14.809396  
  
      obs_ci_upper  
0      34.124758
```

1.2 Part (b) — Plot mpg vs horsepower and the regression line

```
[14]: fig, ax = plt.subplots(figsize=(10,6)) # make figure wider and taller  
ax.scatter(auto["horsepower"], auto["mpg"], s=20, alpha=0.7) # smaller dots, ↴  
        slightly transparent  
ax.plot(auto["horsepower"], model.predict(sm.add_constant(auto["horsepower"])), ↴  
        color='red') # regression line  
ax.set_xlabel("Horsepower")  
ax.set_ylabel("MPG")  
ax.set_title("MPG vs Horsepower with Regression Line")  
plt.show()
```




1.3 Part (c) — Diagnostic Plots

These plots help evaluate assumptions such as linearity, constant variance, and normality of residuals.

```
[16]: import statsmodels.api as sm
import matplotlib.pyplot as plt

fig = plt.figure(figsize=(12, 10)) # make the figure larger
sm.graphics.plot_regress_exog(model, "horsepower", fig=fig)
plt.show()
```



1.3.1 Comments on Diagnostics

- There appears to be some curvature in the residuals, suggesting the relationship may not be perfectly linear.
- There is some evidence of non-constant variance (funnel shape), which means prediction accuracy varies across horsepower values.
- A few points may be potential outliers influencing the model.

Overall, the model shows a clear negative relationship, but improvements such as polynomial regression might yield a better fit.

[]: